

Dimensionality reduction techniques:

Part 1

Sanzhar Askaruly (San)

Ulsan National Institute of Science and Technology
Ph.D. Candidate in Biomedical Engineering

CodeSeoul MLA
December 10, 2022

Overview

1 Motivation

2 Background mathematics

- Population vs sample
- Mean, standard deviation, variance
- Covariance, covariance matrix
- Eigenvectors, eigenvalues

3 Dimensionality reduction techniques

- PCA (Principal component analysis)
 - Example with IRIS dataset
 - PCA under the hood
- t-SNE (Stochastic neighbor embedding)

4 Summary

- Conclusion
- Practicum

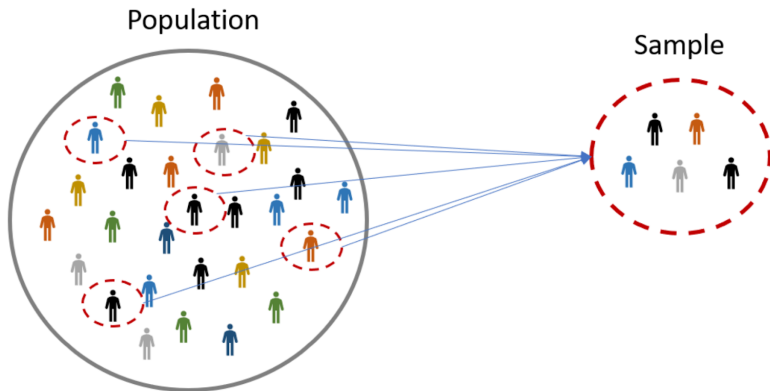
Motivation

Why dimensionality reduction?

Important "toolkit" in machine learning:

- Preprocessing step for other algorithms (reduce size)
- Visualization
- Compression
- Interpolation

Population vs sample



Population vs sample

Population mean

$$\mu = \frac{\sum_{i=1}^N x_i}{N} \quad (1)$$

N is number of items in the population

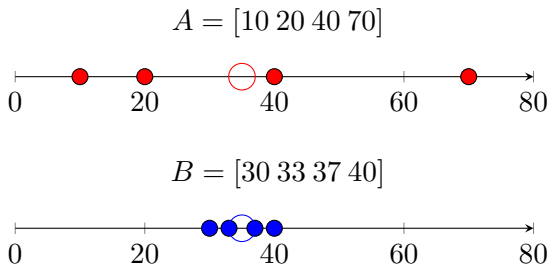
Sample mean

$$\bar{X} = \frac{\sum_{i=1}^n x_i}{n} \quad (2)$$

n is number of items in the sample

Standard deviation

Let's take a look on two samples:



Here, $\bar{A} = \bar{B} = 35$. Unfortunately, mean doesn't tell us a lot except for a middle point.

Standard deviation

For our two sets, $A = [10\ 20\ 40\ 70]$ and $B = [30\ 33\ 37\ 40]$, we would be more interested in the *spread* of the data. So, how do we calculate it?

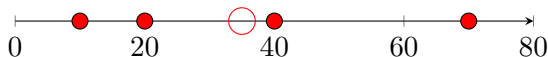
Standard deviation

$$s = \sqrt{\frac{\sum_{i=1}^n (X_i - \bar{X})^2}{(n - 1)}} \quad (3)$$

In plain English, it is the "average distance from the mean of the data set to a point."

Standard deviation

Set 1: $A = [10\ 20\ 40\ 70]$, and $\bar{A} = 35$

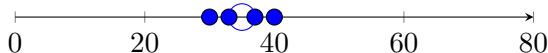


Let's calculate standard deviation:

A	$(A - \bar{A})$	$(A - \bar{A})^2$
10	-25	625
20	-15	225
40	5	25
70	35	1,225
Total		2,100
Divided by (n-1)		700
Square root		26.4575

Standard deviation

Set 2: $B = [30\ 33\ 37\ 40]$, and $\bar{B} = 35$



Let's calculate standard deviation:

B	$(B - \bar{B})$	$(B - \bar{B})^2$
30	-5	25
33	-2	4
37	2	4
40	5	25
Total		58
Divided by (n-1)		19.333
Square root		4.397

Variance

Similar to standard deviation So, how do we calculate it?

Variance

$$s^2 = \frac{\sum_{i=1}^n (X_i - \bar{X})^2}{(n - 1)} \quad (4)$$

Almost identical to the standard deviation (there is no square root in the formula).

Covariance

Previous measures are 1-dimensional. How do we check relationship between the dimensions?

Covariance

Variance

$$\text{var}(X) = \frac{\sum_{i=1}^n (X_i - \bar{X})(X_i - \bar{X})}{(n - 1)} \quad (5)$$

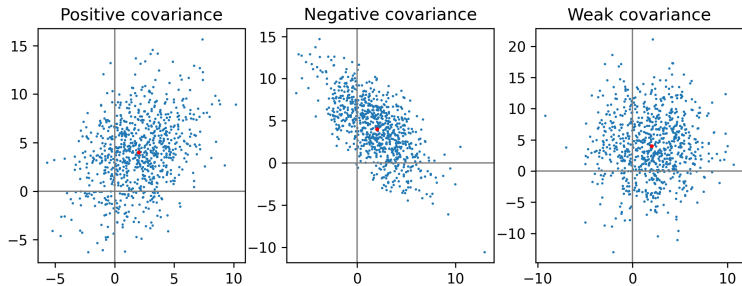
Covariance

$$\text{covar}(X, Y) = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{(n - 1)} \quad (6)$$

How much dimensions vary from the mean *with respect to each other*.

- Covariance between one dimension and itself gives variance.

Covariance



Covariance

Exercise: Find covariance between 2-dimensional dataset

Item number:	1	2	3	4	5
x	10	39	19	23	28
y	43	13	32	21	20

Ans: -120.55

Covariance matrix

When we have more than two dimensions, covariances between individual dimensions could be described in covariance matrix

Covariance matrix

$$C^{n \times n} = (c_{i,j} : c_{i,j} = \text{cov}(\text{Dim}_i, \text{Dim}_j)) \quad (7)$$

where $C^{m \times n}$ is a matrix with n rows and m columns. Example for three dimensions:

$$c = \begin{pmatrix} \text{cov}(x, x) & \text{cov}(x, y) & \text{cov}(x, z) \\ \text{cov}(y, x) & \text{cov}(y, y) & \text{cov}(y, z) \\ \text{cov}(z, x) & \text{cov}(z, y) & \text{cov}(z, z) \end{pmatrix} \quad (8)$$

Since $\text{cov}(a, b) = \text{cov}(b, a)$, the matrix is symmetrical about the main diagonal.


Covariance matrix

Exercise: Find covariance matrix for the 3-dimensional dataset

Item number:	1	2	3
x	1	-1	4
y	2	1	3
z	1	3	-1

Ans:
$$\begin{pmatrix} 4.222 & 1.667 & -3.333 \\ 1.667 & 0.667 & -1.333 \\ -3.333 & -1.333 & 2.667 \end{pmatrix}$$

Eigenvectors



- Eigenvectors (red) do not change direction when a linear transformation (e.g. scaling) is applied to them. Other vectors (yellow) do.

Eigenvectors

An example of non-eigenvector:

$$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix} \times \boxed{\begin{pmatrix} 1 \\ 3 \end{pmatrix}} = \begin{pmatrix} 11 \\ 5 \end{pmatrix} \quad (9)$$

An example of eigenvector:

$$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix} \times \boxed{\begin{pmatrix} 3 \\ 2 \end{pmatrix}} = \begin{pmatrix} 12 \\ 8 \end{pmatrix} = 4 \times \begin{pmatrix} 3 \\ 2 \end{pmatrix} \quad (10)$$

Eigenvalues

The amount by which the original vector was scaled after multiplication by the square matrix

$$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix} \times \begin{pmatrix} 3 \\ 2 \end{pmatrix} = \begin{pmatrix} 12 \\ 8 \end{pmatrix} = 4 \times \begin{pmatrix} 3 \\ 2 \end{pmatrix} \quad (11)$$

What is Principal Component Analysis?

A way of identifying patterns in data: by highlighting their similarities and differences

Pros

- Once patterns in the data, and you compress the data, ie. by reducing the number of dimensions, without much loss of information.

Cons

- Linear \therefore difficult to unfold single ambiguous object really belongs in several disparate locations in the low-dimensional space.

Iris flower dataset

iris setosa



petal

sepal

iris versicolor



petal

sepal

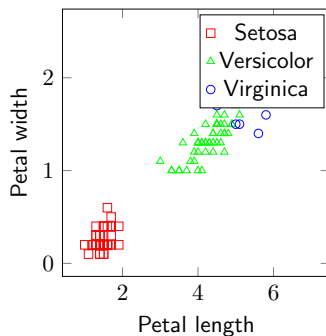
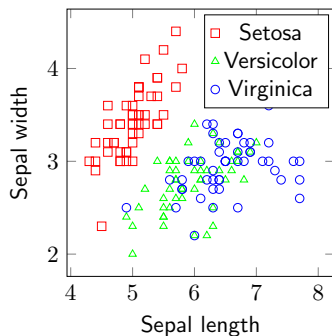
iris virginica



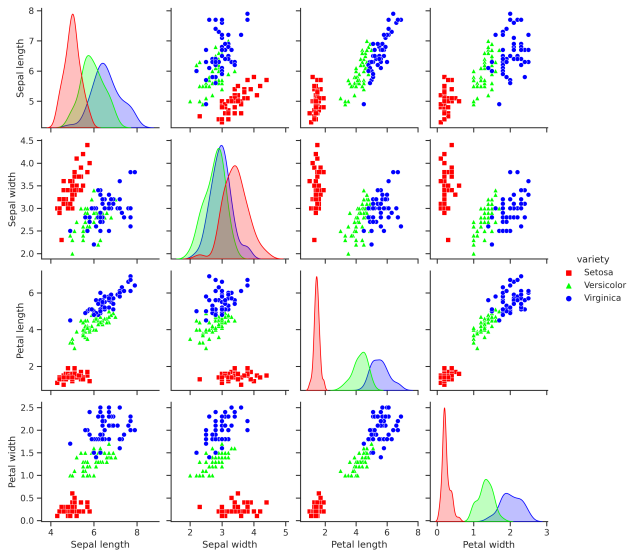
petal

sepal

Iris flower dataset



Iris flower dataset



Applying PCA to Iris flower dataset

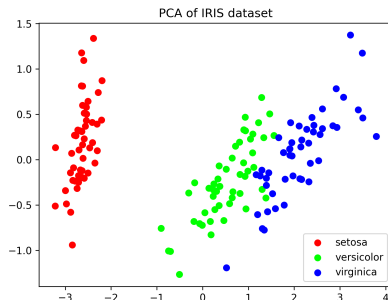
PCA automatically finds principal components (eg. *axes*) in feature space. Here, we defined two, so we can visualize in 2D easily.

```
from sklearn import datasets
from sklearn.decomposition import PCA

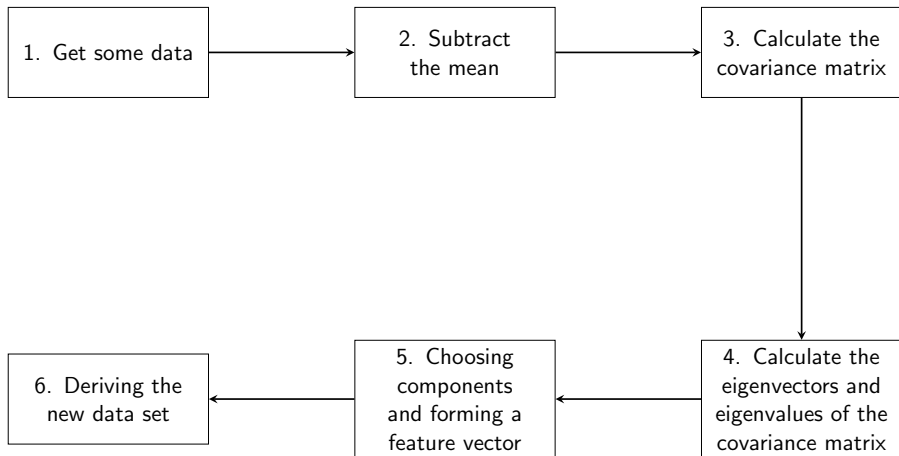
iris = datasets.load_iris()

X = iris.data
y = iris.target
target_names = iris.target_names

pca = PCA(n_components=2)
X_r = pca.fit(X).transform(X)
```



PCA under the hood



PCA under the hood

```
import numpy as np
from numpy.linalg import eig

def PCA_numpy(X,y):
    def mean(x): # np.mean(X, axis = 0)
        return sum(x)/len(x)

    def std(x): # np.std(X, axis = 0)
        return (sum((i - mean(x))**2 for i in x)/len(x))**0.5

    def Standardize_data(X):
        return (X - mean(X))/std(X)

    def covariance(x):
        return (x.T @ x)/(x.shape[0]-1)

    # Step 1: Standardize the data
    X_std = Standardize_data(X)
    # Step 2: Find the covariance matrix
    cov_mat = covariance(X_std) # np.cov(X_std.T)

    # Step 3: Find the eigenvectors and eigenvalues of the covariance matrix
    eig_vals, eig_vecs = eig(cov_mat)
```

PCA under the hood

```
max_abs_idx = np.argmax(np.abs(eig_vecs), axis=0)
signs = np.sign(eig_vecs[max_abs_idx, range(eig_vecs.shape[0])])
eig_vecs = eig_vecs*signs[np.newaxis,:]
```

```
# Step 4: Rearrange the eigenvectors and eigenvalues
```

```
eig_pairs = [(np.abs(eig_vals[i]), eig_vecs[i,:]) for i in range(len(eig_vals))]
```

```
# Then, we sort the tuples from the highest to the lowest based on eigenvalues magni
```

```
eig_pairs.sort(key=lambda x: x[0], reverse=True)
```

```
eig_vals_sorted = np.array([x[0] for x in eig_pairs])
```

```
eig_vecs_sorted = np.array([x[1] for x in eig_pairs])
```

```
# Step 5: Choose principal components
```

```
k = 2
```


```
W = eig_vecs_sorted[:k, :] # Projection matrix
```

```
# Step 6: Project the data
```

```
X_proj = X_std.dot(W.T)
```

```
return X_proj
```

PCA under the hood



t-SNE (Stochastic neighbor embedding)

To be continued

Conclusion

Thank you for your attention!

References